

Handling-missing-values.R

patriciamaya

2020-12-06

#HANDLING MISSING VALUES

```
data <- read.csv("~/Downloads/hmeq.csv")
summary(data)
```

```
##      BAD      LOAN      MORTDUE      VALUE
## Min.   :0.0000  Min.   : 1100  Min.   : 2063  Min.   : 8000
## 1st Qu.:0.0000  1st Qu.:11100  1st Qu.: 46276  1st Qu.: 66076
## Median :0.0000  Median :16300  Median : 65019  Median : 89236
## Mean   :0.1995  Mean   :18608  Mean   : 73761  Mean   :101776
## 3rd Qu.:0.0000  3rd Qu.:23300  3rd Qu.: 91488  3rd Qu.:119824
## Max.   :1.0000  Max.   :89900  Max.   :399550  Max.   :855909
##                                     NA's   :518   NA's   :112
##      REASON      JOB      YOJ      DEROG
## Length:5960     Length:5960     Min.   : 0.000  Min.   : 0.0000
## Class :character Class :character  1st Qu.: 3.000  1st Qu.: 0.0000
## Mode  :character Mode  :character  Median : 7.000  Median : 0.0000
##                                     Mean   : 8.922  Mean   : 0.2546
##                                     3rd Qu.:13.000  3rd Qu.: 0.0000
##                                     Max.   :41.000  Max.   :10.0000
##                                     NA's   :515   NA's   :708
##      DELINQ      CLAGE      NINQ      CLNO
## Min.   : 0.0000  Min.   : 0.0  Min.   : 0.000  Min.   : 0.0
## 1st Qu.: 0.0000  1st Qu.:115.1  1st Qu.: 0.000  1st Qu.:15.0
## Median : 0.0000  Median :173.5  Median : 1.000  Median :20.0
## Mean   : 0.4494  Mean   :179.8  Mean   : 1.186  Mean   :21.3
## 3rd Qu.: 0.0000  3rd Qu.:231.6  3rd Qu.: 2.000  3rd Qu.:26.0
## Max.   :15.0000  Max.   :1168.2  Max.   :17.000  Max.   :71.0
## NA's   :580     NA's   :308   NA's   :510   NA's   :222
##      DEBTINC
## Min.   : 0.5245
## 1st Qu.: 29.1400
## Median : 34.8183
## Mean   : 33.7799
## 3rd Qu.: 39.0031
## Max.   :203.3121
## NA's   :1267
```

Handling missing values of NUMERICAL variables

```
sum(is.na(data$NINQ)) # This give you the number of missing values in the variable NINQ
```

```
## [1] 510
```

```
sum(complete.cases(data$NINQ)) # Count of complete cases in the variable NINQ
```

```
## [1] 5450
```

```
sum(!complete.cases(data$NINQ)) # Count of NOT complete cases in the variable NINQ
```

```
## [1] 510
```

```
which(!complete.cases(data$NINQ)) # Which cases (row numbers) are NOT complete
```

```
## [1] 4 11 18 52 64 74 96 106 113 116 128 140 144 145 146
## [16] 153 155 160 165 166 170 172 174 187 191 212 218 222 227 230
## [31] 232 238 240 242 246 252 266 269 285 293 300 303 305 310 318
## [46] 323 331 334 337 339 343 344 347 351 353 357 358 359 362 366
## [61] 367 375 381 382 390 392 396 397 399 402 414 418 421 425 432
## [76] 435 444 465 469 473 482 490 503 527 532 536 537 545 550 561
## [91] 566 567 597 601 604 609 619 634 643 645 649 654 669 688 692
## [106] 703 711 717 726 735 737 748 749 752 763 765 770 772 788 783
## [121] 786 790 812 818 830 844 854 858 865 868 882 899 922 929 932
## [136] 933 935 947 970 974 975 980 987 988 992 1011 1031 1040 1048 1076
## [151] 1084 1092 1094 1106 1123 1138 1145 1146 1153 1155 1157 1182 1196 1210 1224
## [166] 1236 1238 1244 1249 1254 1257 1276 1296 1321 1333 1336 1339 1348 1361 1364
## [181] 1373 1389 1395 1402 1406 1411 1417 1422 1426 1427 1434 1468 1488 1499 1505
## [196] 1508 1532 1554 1556 1568 1571 1573 1589 1592 1628 1635 1645 1646 1663 1675
## [211] 1688 1690 1695 1736 1766 1775 1788 1790 1824 1842 1864 1878 1895 1898 1960
## [226] 1961 1965 1967 1989 1998 2062 2067 2073 2102 2108 2113 2121 2127 2154 2166
## [241] 2168 2218 2244 2246 2266 2267 2297 2304 2309 2310 2342 2355 2357 2379 2397
## [256] 2412 2417 2427 2440 2450 2464 2473 2476 2515 2518 2543 2551 2588 2626 2675
## [271] 2680 2686 2690 2743 2752 2759 2814 2876 2891 2906 2936 2970 2984 2988 2993
## [286] 2998 3006 3047 3051 3062 3077 3097 3125 3135 3136 3141 3152 3157 3177 3250
## [301] 3308 3351 3366 3385 3392 3478 3491 3518 3556 3582 3601 3621 3623 3624 3631
```

```
## [316] 3652 3673 3683 3695 3697 3706 3721 3737 3744 3746 3747 3762 3768 376
9 3777
## [331] 3817 3822 3828 3835 3840 3842 3844 3856 3860 3874 3886 3918 3937 394
3 3953
## [346] 3955 3978 3993 4003 4013 4017 4036 4072 4073 4074 4097 4113 4119 412
2 4123
## [361] 4128 4131 4142 4157 4160 4162 4178 4200 4207 4211 4230 4242 4252 425
8 4273
## [376] 4274 4279 4281 4295 4303 4309 4321 4322 4323 4325 4343 4344 4351 435
2 4361
## [391] 4364 4366 4382 4385 4387 4393 4402 4414 4419 4421 4425 4430 4432 443
8 4439
## [406] 4476 4480 4482 4500 4542 4544 4549 4567 4573 4575 4576 4582 4585 458
9 4600
## [421] 4610 4633 4645 4657 4661 4671 4672 4681 4683 4698 4716 4748 4759 478
4 4785
## [436] 4786 4790 4791 4795 4801 4803 4818 4846 4851 4852 4858 4866 4868 487
0 4881
## [451] 4900 4910 4932 4942 4946 4948 4973 4975 4976 4989 4991 5000 5045 504
9 5061
## [466] 5072 5076 5102 5107 5113 5127 5148 5165 5200 5229 5230 5248 5253 526
5 5274
## [481] 5275 5314 5341 5347 5366 5386 5418 5423 5434 5452 5464 5467 5472 549
2 5496
## [496] 5699 5701 5702 5719 5746 5749 5751 5752 5755 5757 5760 5763 5766 580
9 5838
```

```
# The function "na.omit()" DELETES ALL instances with missing values and retu
rns
# the object with listwise deletion of missing values.
```


```
NINQ_Imputed = na.omit(data$NINQ) # Create new variable without missing value
s
sum(is.na(NINQ_Imputed))
```

```
## [1] 0
```

```
# REPLACE missing values by a particular value (mean)
data$NINQ[is.na(data$NINQ)] = mean(data$NINQ, na.rm=TRUE) # Recode all NA in
NINQ as the average value
sum(is.na(data$NINQ))
```

```
## [1] 0
```

```
# REPLACING using mice for looking at missing DATA PATTERN
library(mice)
md.pattern(data)
```



##	BAD	LOAN	REASON	JOB	NINQ	VALUE	CLNO	CLAGE	YOJ	MORTDUE	DELINQ	DEROG	DEBTINC
## 3551	1	1	1	1	1	1	1	1	1	1	1	1	1
## 932	1	1	1	1	1	1	1	1	1	1	1	1	0
## 158	1	1	1	1	1	1	1	1	1	1	1	0	1
## 33	1	1	1	1	1	1	1	1	1	1	1	0	0
## 29	1	1	1	1	1	1	1	1	1	1	0	1	1
## 8	1	1	1	1	1	1	1	1	1	1	0	1	0
## 178	1	1	1	1	1	1	1	1	1	1	0	0	1
## 22	1	1	1	1	1	1	1	1	1	1	0	0	0
## 176	1	1	1	1	1	1	1	1	1	0	1	1	1
## 55	1	1	1	1	1	1	1	1	1	0	1	1	0
## 12	1	1	1	1	1	1	1	1	1	0	0	1	1
## 1	1	1	1	1	1	1	1	1	1	0	0	1	0
## 23	1	1	1	1	1	1	1	1	1	0	0	0	1
## 2	1	1	1	1	1	1	1	1	1	0	0	0	0
## 188	1	1	1	1	1	1	1	1	0	1	1	1	1
## 40	1	1	1	1	1	1	1	1	0	1	1	1	0
## 1	1	1	1	1	1	1	1	1	0	1	1	0	0
## 10	1	1	1	1	1	1	1	1	0	1	0	1	1
## 2	1	1	1	1	1	1	1	1	0	1	0	1	0
## 7	1	1	1	1	1	1	1	1	0	1	0	0	1
## 1	1	1	1	1	1	1	1	1	0	1	0	0	0
## 53	1	1	1	1	1	1	1	1	0	0	1	1	1
## 13	1	1	1	1	1	1	1	1	0	0	1	1	0
## 6	1	1	1	1	1	1	1	1	0	0	1	0	1
## 1	1	1	1	1	1	1	1	1	0	0	1	0	0
## 12	1	1	1	1	1	1	1	1	0	0	0	1	1
## 1	1	1	1	1	1	1	1	1	0	0	0	1	0
## 40	1	1	1	1	1	1	1	1	0	0	0	0	1
## 4	1	1	1	1	1	1	1	1	0	0	0	0	0
## 28	1	1	1	1	1	1	1	0	1	1	1	1	1
## 17	1	1	1	1	1	1	1	0	1	1	1	1	0

[illegible]

```
## 22      5
## 4       5
## 9       6
## 43      5
## 8       6
## 6       3
## 1       4
## 3       5
## 1       2
## 1       3
## 1       4
## 2       5
## 4       6
## 2       7
## 6       7
## 4       8
##      4230
```

The output tells us that 3551 samples are complete, 932 samples miss only DEBTINC, 158 samples miss only the DEROG and so on.

The “mice()” function takes care of imputing process.

```
NewData = mice(data, m=5, maxit=50, meth="pmm", seed=500)
```

```
summary(NewData)
```

```
## Class: mids
## Number of multiple imputations: 5
## Imputation methods:
##      BAD      LOAN MORTDUE  VALUE  REASON      JOB      YOJ      DEROG  DELINQ  CLAGE
##      ""      ""      "pmm"  "pmm"      ""      ""      "pmm"  "pmm"  "pmm"  "pmm"
##      NINQ     CLNO DEBTINC
##      ""      "pmm"  "pmm"
## PredictorMatrix:
##      BAD LOAN MORTDUE VALUE REASON JOB YOJ DEROG DELINQ CLAGE NINQ CLNO
## BAD      0    1      1    1      0  0  1    1    1    1    1    1
## LOAN      1    0      1    1      0  0  1    1    1    1    1    1
## MORTDUE    1    1      0    1      0  0  1    1    1    1    1    1
## VALUE      1    1      1    0      0  0  1    1    1    1    1    1
## REASON     1    1      1    1      0  0  1    1    1    1    1    1
## JOB        1    1      1    1      0  0  1    1    1    1    1    1
##      DEBTINC
## BAD      1
## LOAN      1
## MORTDUE    1
## VALUE      1
## REASON     1
## JOB        1
## Number of logged events: 2
##   it im dep      meth  out
## 1  0  0      constant REASON
## 2  0  0      constant  JOB
```

m=5 refers to the number of imputed datasets. Five is the default value

meth='pmm' refers to the imputation method.

#In this case we are using predictive mean matching as imputation method

We can get back the completed dataset using the complete() function

```
New_data <- as.data.frame(complete(NewData, 1))
head(New_data)
```

```
##      BAD LOAN MORTDUE  VALUE  REASON      JOB  YOJ  DEROG  DELINQ      CLAGE      N
INQ
## 1    1 1100    25860  39025 HomeImp  Other 10.5      0      0  94.36667 1.000
000
## 2    1 1300    70053  68400 HomeImp  Other  7.0      0      2 121.83333 0.000
000
## 3    1 1500    13500  16700 HomeImp  Other  4.0      0      0 149.46667 1.000
000
## 4    1 1500    72136  85100                22.0      0      0  62.35974 1.186
055
## 5    0 1700    97800 112000 HomeImp  Office 3.0      0      0  93.33333 0.000
000
## 6    1 1700    30548  40320 HomeImp  Other  9.0      0      0 101.46600 1.000
000
##      CLNO  DEBTINC
## 1        9 30.60700
## 2       14 42.58162
## 3       10 37.33952
## 4       15 36.43872
## 5       14 29.29518
## 6        8 37.11361
```

#Handling missing values of CATEGORICAL variables:

#As far as categorical variables are concerned, replacing categorical variables is usually not advisable. Some common practice include replacing missing categorical variables with the mode of the observed ones, however, it is questionable whether it is a good choice.

```
data$REASON <- as.factor(data$REASON)
data$JOB <- as.factor(data$JOB)
```

#remove NA values if only instances of categorical variables are missing
Data <- data[complete.cases(data),] *#removes ALL rows where a value is missing*

#we need to first deal with numerical missing values and then categorical variables.