

## LOAN APPROVAL USING PREDICTIVE ANALYTICS

To minimize loss from the bank's perspective, the bank needs a decision rule regarding who to give approval of the loan and who not to. An applicant's demographic and socio-economic profiles are considered by loan managers before a decision is taken regarding his/her loan application. The Bank.data contains data on 9 input variables and the classification target indicating whether an applicant is considered a **Good or a Bad credit risk** for 1000 loan applicants. A predictive model developed on this data is expected to provide a bank manager guidance for making a decision whether to approve a loan to a prospective applicant based on his/her profiles.

The variables in this dataset are:

- Age (numeric)
- Sex (categorical: male, female)
- Job (categorical: 0 - unskilled and non-resident, 1 - unskilled and resident, 2 - skilled, 3 - highly skilled)
- Housing (categorical: own, rent, or free)
- Saving accounts (categorical: little, moderate, quite rich, rich)
- Checking account (categorical: little, moderate, rich)
- Credit amount (numeric, in USD)
- Duration (numeric, in month)
- Purpose (categorical: car, furniture/equipment, radio/TV, domestic appliances, repairs, education, business, vacation/others)
- Target (categorical: 1 - Good, 0 - Bad)

```
data <- read.csv("~/Downloads/Bank.data.csv")
head(data)
```

```
##   ID Age  Sex Job Housing Saving.accounts Checking.account Credit.amount
## 1  0  67 male  2   own              <NA>         little        1169
## 2  1  22 female 2   own          little         moderate        5951
## 3  2  49 male  1   own          little              <NA>        2096
## 4  3  45 male  2   free          little         little        7882
## 5  4  53 male  2   free          little         little        4870
## 6  5  35 male  1   free              <NA>              <NA>        9055
##   Duration              Purpose Target
## 1         6          radio/TV       1
## 2        48          radio/TV       0
## 3        12          education       1
## 4        42 furniture/equipment       1
## 5        24              car       0
## 6        36          education       1
```

```
#formatting data
```

```
data$Sex <- as.factor(data$Sex)
data$Job <- as.factor(data$Job)
data$Housing <- as.factor(data$Housing)
data$Saving.accounts <- as.factor(data$Saving.accounts)
data$Checking.account <- as.factor(data$Checking.account)
data$Purpose <- as.factor(data$Purpose)
data$Target <- as.factor(data$Target)
```

```
data$Age <- as.numeric(data$Age)
data$Credit.amount <- as.numeric(data$Credit.amount)
data$Duration <- as.numeric(data$Duration)
```

```
summary(data)
```

```
##      ID      Age      Sex      Job      Housing
## Min.   : 0.0   Min.   :19.00 female:310  0: 22   free:108
## 1st Qu.:249.8 1st Qu.:27.00 male  :690  1:200   own :713
## Median :499.5 Median :33.00           2:630   rent:179
## Mean   :499.5 Mean   :35.55           3:148
## 3rd Qu.:749.2 3rd Qu.:42.00
## Max.   :999.0 Max.   :75.00
##
##      Saving.accounts Checking.account Credit.amount      Duration
## little      :603      little :274      Min.   : 250      Min.   : 4.0
## moderate    :103      moderate:269      1st Qu.: 1366      1st Qu.:12.0
## quite rich: 63      rich      : 63      Median : 2320      Median :18.0
## rich        : 48      NA's      :394      Mean   : 3271      Mean   :20.9
## NA's        :183                        3rd Qu.: 3972      3rd Qu.:24.0
##                                         Max.   :18424      Max.   :72.0
##
##
##      Purpose      Target
## car              :337    0:300
## radio/TV          :280    1:700
## furniture/equipment:181
## business          : 97
## education         : 59
## repairs           : 22
## (Other)           : 24
```

```
#5 number summary for the credit amount (minimum, 1st quartile, median, 3rd quartile, maximum)
```

```
summary(data$Credit.amount)
```


```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      250   1366   2320   3271   3972   18424
```

```
#correlation of Duration and Credit
```

```
cor(data$Duration, data$Credit.amount)
```

```
## [1] 0.6249842
```

```
plot(data$Duration, data$Credit.amount)
```




Variables are NOT highly correlated so we can keep both variables for the analysis (Usually, high correlation is considered above 0.75)

*# Distribution of Credit amount for applicants considered as “Good” and “Bad” credit risk.*

```
boxplot(data$Credit.amount ~ data$Target, main = "Distribution of Credit amount", xlab = "0 - Bad", ylab = "Credit Amount", title = "xccc", col = "lightblue")
```

**Distribution of Credit amount**



```
# table that contains the frequency of different housing types (free, own, rent) for "Good" and "Bad" instances.
table(data$Housing, data$Target, useNA = "ifany", dnn= c("Housing", "Credit"))
```

```
##           Credit
## Housing    0    1
##   free   44   64
##   own  186  527
##   rent   70  109
```

```
#Renaming some variables
```

```
names(data)[names(data) == "Checking.account"] <- "Checkings"
names(data)[names(data) == "Saving.accounts"] <- "Savings"
names(data)[names(data) == "Credit.amount"] <- "Credit"
```

## HANDLING MISSING VALUES

```
#Handling missing values
```

```
library(mice)
```

```
## Warning: package 'mice' was built under R version 4.0.2
```

```
##
```


```
## Attaching package: 'mice'
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      cbind, rbind
```

```
md.pattern(data)
```



```
##           ID Age Sex Job Housing Credit Duration Purpose Target Savings Checking
## 522  1    1    1    1          1        1          1        1        1        1
## 295  1    1    1    1          1        1          1        1        1        1
## 84   1    1    1    1          1        1          1        1        1        1
## 99   1    1    1    1          1        1          1        1        1        2
```

```
## 84  1  1  1  1      1      1      1      1      1      0
1  1
## 99  1  1  1  1      1      1      1      1      1      0
0  2
##      0  0  0  0      0      0      0      0      0      183      39
4 577
```

The output tells us that 522 samples are complete, 295 samples miss ONLY Checking.account, 84 samples miss only the saving.accounts and 99 samples miss both Saving.accounts and Checking.account.

*#As far as categorical variables are concerned, replacing categorical variables is usually not advisable. Some common practice include replacing missing categorical variables with the mode of the observed ones, however, it is questionable whether it is a good choice.*

*#Since all NA values are from categorical variables:*

*#REMOVE ALL ROWS WITH NA VALUES*

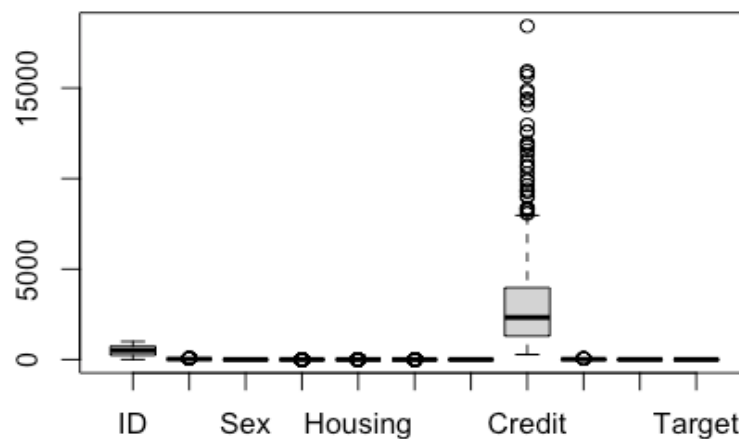
```
Data <- na.omit(data)
sum(is.na(Data))
```

```
## [1] 0
```

## HANDLING OUTLIERS

```
library(dplyr)
```

```
boxplot(Data)
```



*#credit variable*

```
credit_outliers = boxplot.stats(Data$Credit)$out # We first save all the outliers in the vector
credit_outliers
```

```
## [1] 12579 14421 12612 15945 11938 10623 10961 8978 11998 10722 9398 9960
```

```
## [13] 14782 14318 12976 11760 11328 8318 9034 8086 9857 14027 11560 80
65
## [25] 9271 9283 9629 15857 11816 15672 18424 14896 10297 8358 8386 82
29
```

*#duration variable*

Duration\_outliers = `boxplot.stats`(Data\$Duration)\$out *# We first save all the outliers in the vector*

Duration\_outliers

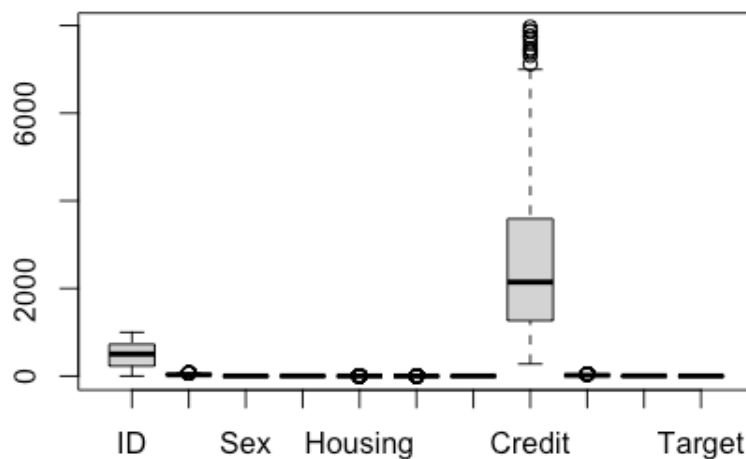
```
## [1] 60 54 60 60 72 60 60 60
```

*#REMOVING OUTLIERS FROM DATA*

Data<- Data[`-which`(Data\$Duration `%in%` Duration\_outliers),]

Data<- Data[`-which`(Data\$Credit `%in%` credit\_outliers),]

`boxplot`(Data)



## Train/Test split

```
set.seed(123)
```

```
indx <- sample(2,nrow(Data), replace=TRUE, prob = c(0.7, 0.3))
```

```
train <- Data[indx==1, ]
```

```
test <- Data[indx==2, ]
```

## Logistic Regression Model using Forward Selection Technique

Forward Selection considers one variable at a time:

-If the variable improves the model (reduces AIC), we include it

-Otherwise, we don't include it.

We look at r-squared, f-test or AIC.

We will check all possibilities from the null case to the full case

### #Extreme Cases

```
full <- glm(Target ~ . , data=Data, family = "binomial")
null <- glm(Target ~ 1 , data=Data, family = "binomial")
step(null, scope = list(lower=null, upper=full), direction="forward")
```

## Start: AIC=652.98

## Target ~ 1

##

	Df	Deviance	AIC
## + Duration	1	622.77	626.77
## + Checkings	2	638.19	644.19
## + Savings	3	640.57	648.57
## + Housing	2	643.20	649.20
## + Sex	1	648.04	652.04
## + Age	1	648.46	652.46
## + ID	1	648.80	652.80
## <none>		650.98	652.98
## + Credit	1	650.83	654.83
## + Job	3	649.00	657.00
## + Purpose	7	641.08	657.08

##

## Step: AIC=626.77

## Target ~ Duration

##

	Df	Deviance	AIC
## + Credit	1	609.26	615.26
## + Checkings	2	612.15	620.15
## + Sex	1	617.87	623.87
## + Savings	3	614.97	624.97
## + Housing	2	617.31	625.31
## + ID	1	619.91	625.91
## + Age	1	620.16	626.16
## <none>		622.77	626.77
## + Purpose	7	608.85	626.85
## + Job	3	620.28	630.28

##

## Step: AIC=615.26

## Target ~ Duration + Credit

##

	Df	Deviance	AIC
## + Checkings	2	597.86	607.86
## + Housing	2	601.96	611.96
## + Savings	3	600.45	612.45
## + Sex	1	605.29	613.29
## + Purpose	7	594.51	614.51
## + Age	1	607.10	615.10
## <none>		609.26	615.26
## + ID	1	607.27	615.27
## + Job	3	608.57	620.57

##

```

## Step: AIC=607.86
## Target ~ Duration + Credit + Checkings
##
##           Df Deviance    AIC
## + Sex      1   593.64 605.64
## + Housing   2   591.77 605.77
## + Savings   3   590.55 606.55
## <none>      597.86 607.86
## + Age       1   595.87 607.87
## + Purpose   7   583.96 607.96
## + ID        1   596.39 608.39
## + Job       3   597.12 613.12
##
## Step: AIC=605.64
## Target ~ Duration + Credit + Checkings + Sex
##
##           Df Deviance    AIC
## + Savings   3   585.88 603.88
## + Housing   2   589.07 605.07
## <none>      593.64 605.64
## + Purpose   7   579.86 605.86
## + Age       1   592.33 606.33
## + ID        1   592.52 606.52
## + Job       3   593.15 611.15
##
## Step: AIC=603.88
## Target ~ Duration + Credit + Checkings + Sex + Savings
##
##           Df Deviance    AIC
## + Housing   2   580.99 602.99
## + Purpose   7   571.70 603.70
## <none>      585.88 603.88
## + Age       1   584.86 604.86
## + ID        1   584.97 604.97
## + Job       3   585.07 609.07
##
## Step: AIC=602.99
## Target ~ Duration + Credit + Checkings + Sex + Savings + Housing
##
##           Df Deviance    AIC
## <none>      580.99 602.99
## + Age       1   579.73 603.73
## + ID        1   580.19 604.19
## + Purpose   7   568.19 604.19
## + Job       3   579.85 607.85
##
## Call: glm(formula = Target ~ Duration + Credit + Checkings + Sex +
##           Savings + Housing, family = "binomial", data = Data)
##

```



```
## Coefficients:
##      (Intercept)      Duration      Credit Checkingsmoderate
##      0.2091507      -0.0728745      0.0002961      0.2653101
##      Checkingsrich      Sexmale      Savingsmoderate Savingsquite rich
##      1.0498166      0.3809951      0.0592703      0.5239743
##      Savingsrich      Housingown      Housingrent
##      1.5387266      0.4530058      -0.0441255
##
## Degrees of Freedom: 480 Total (i.e. Null); 470 Residual
## Null Deviance:      651
## Residual Deviance: 581 AIC: 603
```

$AIC = -2\log(\text{Likelihood}) + 2K$

$AIC = \text{Residual deviance} + 2 \times \text{number of parameters.}$

AIC is a single number score that can be used to determine which of multiple models is most likely to be the best model for a given dataset.

A lower AIC score is better.

“The Akaike information criterion (AIC) is an estimator of out-of-sample prediction error and thereby relative quality of statistical models for a given set of data. Given a collection of models for the data, AIC estimates the quality of each model, relative to each of the other models. Thus, AIC provides a means for model selection.”

*#Log Regression with variables selected*

```
Log_Reg <- glm(Target ~ Duration + Credit + Checkings + Sex + Savings + Housing, family = "binomial", data = Data)
```

```
predictions <- predict(Log_Reg, newdata = test, type="response")
```

*#probability of being in class GOOD*

```
Class <- ifelse(predictions >= 0.5 , 1, 0)
```

*#Confusion Matrix*

```
table(test$Target , Class, dnn=c("Predictions", "Actual"))
```

```
##           Actual
## Predictions  0  1
##           0 23 26
##           1 17 70
```

*#accuracy function*

```
accuracy<- function(actual,predictions)
{
  y <- as.vector(table(predictions,actual))
  names(y) <- c("TN","FP","FN","TP")
  accuracy <- (y["TN"] + y["TP"])/ sum(y)
  return(as.numeric(accuracy))
}
```

```

accuracy(test$Target, Class)

## [1] 0.6838235

summary(Log_Reg)

##
## Call:
## glm(formula = Target ~ Duration + Credit + Checkings + Sex +
##      Savings + Housing, family = "binomial", data = Data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.1902  -1.1094   0.6413   0.9470   1.6766
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    2.092e-01  4.278e-01   0.489  0.624885
## Duration      -7.287e-02  1.250e-02  -5.828  5.62e-09 ***
## Credit         2.961e-04  7.829e-05   3.782  0.000155 ***
## Checkingsmoderate 2.653e-01  2.187e-01   1.213  0.225134
## Checkingsrich   1.050e+00  3.693e-01   2.842  0.004478 **
## Sexmale         3.810e-01  2.182e-01   1.746  0.080752 .
## Savingsmoderate  5.927e-02  3.205e-01   0.185  0.853304
## Savingsquite rich 5.240e-01  5.060e-01   1.036  0.300405
## Savingsrich     1.539e+00  6.568e-01   2.343  0.019149 *
## Housingown      4.530e-01  3.322e-01   1.364  0.172637
## Housingrent     -4.413e-02  3.787e-01  -0.117  0.907254
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 650.98  on 480  degrees of freedom
## Residual deviance: 580.99  on 470  degrees of freedom
## AIC: 602.99
##
## Number of Fisher Scoring iterations: 4

```

We can see the significant variables with a \*.

## DECISION TREE

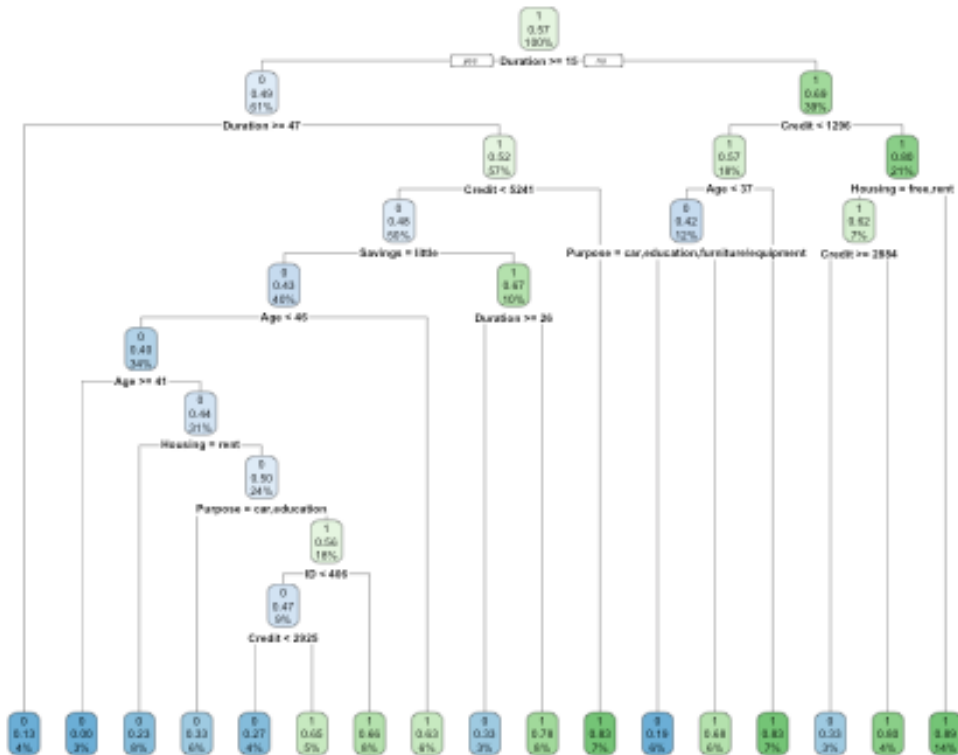
```

library(rpart)

library(rpart.plot)

```

```
tree_model <- rpart(Target ~ . , data=train)
rpart.plot(tree_model)
```



`rpart(formula, data= train, parms= , control= )`

`control` - controls how to split. `control = rpart.control(minsplit=10)`

`minsplit = 10` -> at least 10 instances must be in each node so that it could be split further

`minbucket = 10` -> min num of instances expected in terminal nodes

`cp` - complexity parameter -> is used to control the size of the decision tree and to select the optimal tree size.

want the tree with the min error & also min size of tree

- when `cp` is large – size of tree is small and error is larger
- when `cp` is small – size of tree is large and error is smaller

Example: `rpart(y~., data, parms=list(split=c("information","gini")), cp = 0.01, minsplit=20, minbucket=7, maxdepth=30)`

```
print(tree_model)
```

```

## n= 345
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
##      1) root 345 148 1 (0.4289855 0.5710145)
##          2) Duration>=14.5 211 104 0 (0.5071090 0.4928910)
##              4) Duration>=46.5 15    2 0 (0.8666667 0.1333333) *
##              5) Duration< 46.5 196  94 1 (0.4795918 0.5204082)
##                  10) Credit< 5240.5 173  83 0 (0.5202312 0.4797688)
##                      20) Savings=little 137  59 0 (0.5693431 0.4306569)
##                          40) Age< 45.5 118  47 0 (0.6016949 0.3983051)
##                              80) Age>=40.5 10    0 0 (1.0000000 0.0000000) *
##                              81) Age< 40.5 108  47 0 (0.5648148 0.4351852)
##                                  162) Housing=rent 26    6 0 (0.7692308 0.2307692) *
##                                  163) Housing=free,own 82  41 0 (0.5000000 0.5000000)
##                                      326) Purpose=car,education 21    7 0 (0.6666667 0.3333333)
##                                          *
##                                              327) Purpose=business,furniture/equipment,radio/TV,repairs,
##                                              vacation/others 61  27 1 (0.4426230 0.5573770)
##                                                  654) ID< 406 32  15 0 (0.5312500 0.4687500)
##                                                  1308) Credit< 2925 15    4 0 (0.7333333 0.2666667) *
##                                                  1309) Credit>=2925 17    6 1 (0.3529412 0.6470588) *
##                                                  655) ID>=406 29  10 1 (0.3448276 0.6551724) *
##                                                      41) Age>=45.5 19    7 1 (0.3684211 0.6315789) *
##                                                          21) Savings=moderate,quite rich,rich 36  12 1 (0.3333333 0.66666
##                                                          67)
##                                                              42) Duration>=25.5 9    3 0 (0.6666667 0.3333333) *
##                                                              43) Duration< 25.5 27    6 1 (0.2222222 0.7777778) *
##                                                                  11) Credit>=5240.5 23    4 1 (0.1739130 0.8260870) *
##                                                                      3) Duration< 14.5 134  41 1 (0.3059701 0.6940299)
##                                                                          6) Credit< 1296 63  27 1 (0.4285714 0.5714286)
##                                                                              12) Age< 36.5 40  17 0 (0.5750000 0.4250000)
##                                                                                  24) Purpose=car,education,furniture/equipment 21    4 0 (0.809523
##                                                                                  8 0.1904762) *
##                                                                                        25) Purpose=business,domestic appliances,radio/TV,repairs 19    6
##                                                                                        1 (0.3157895 0.6842105) *
##                                                                                            13) Age>=36.5 23    4 1 (0.1739130 0.8260870) *
##                                                                                                7) Credit>=1296 71  14 1 (0.1971831 0.8028169)
##                                                                                                    14) Housing=free,rent 24    9 1 (0.3750000 0.6250000)
##                                                                                                        28) Credit>=2884 9    3 0 (0.6666667 0.3333333) *
##                                                                                                            29) Credit< 2884 15    3 1 (0.2000000 0.8000000) *
##                                                                                                                15) Housing=own 47    5 1 (0.1063830 0.8936170) *

```

node), split, n, loss, yval, (yprob)

split- variable name & condition

n - number of instances on the node

loss - instances predicted in the wrong class

yval - predicted class

yprob - probability of being in each class [(x,y) first num x corresponds to the class predicted in the first node]

asteriks - terminal nodes

```
tree_pred_class <- predict(tree_model, test, type = "class")
table(tree_pred_class, test$Target)

##
## tree_pred_class  0  1
##                0 18 26
##                1 31 61


accuracy(tree_pred_class, test$Target)

## [1] 0.5808824

#To find best value of CP we can use printcp(), plotcp(), or summary() and choose CP with minimum xerror
printcp(tree_model)

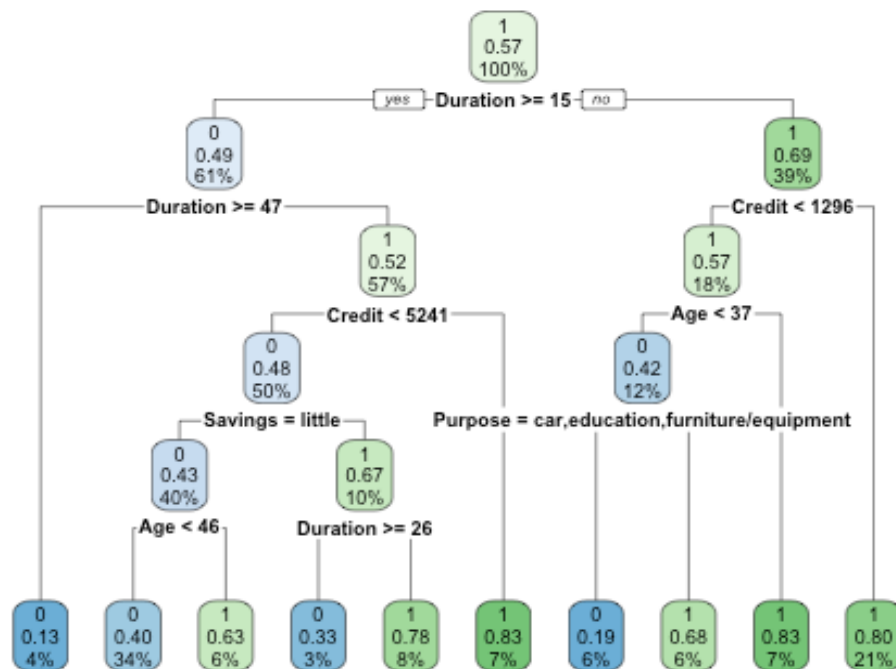
##
## Classification tree:
## rpart(formula = Target ~ ., data = train)
##
## Variables actually used in tree construction:
## [1] Age      Credit  Duration Housing  ID      Purpose  Savings
##
## Root node error: 148/345 = 0.42899
##
## n= 345
##
##      CP nsplit rel error  xerror    xstd
## 1 0.050676      0  1.00000 1.00000 0.062114
## 2 0.033784      4  0.79730 1.06757 0.062529
## 3 0.029279      5  0.76351 1.05405 0.062463
## 4 0.020270      8  0.67568 1.03378 0.062348
## 5 0.015766      9  0.65541 0.97973 0.061948
## 6 0.010135     14  0.56081 1.01351 0.062215
## 7 0.010000     16  0.54054 1.00000 0.062114

plotcp(tree_model)
```



#### Tree with best CP value and minimum prediction error

```
optimal <- which.min(tree_model$cptable[, "xerror"])
cp <- tree_model$cptable[optimal, "CP"]
tree_pruned <- prune(tree_model, cp = cp)
rpart.plot(tree_pruned)
```



```

tree_pruned_pred_class <- predict(tree_pruned, test, type = "class")
table(tree_pruned_pred_class, test$Target)

##
## tree_pruned_pred_class  0  1
##                        0 24 30
##                        1 25 57

accuracy(tree_pruned_pred_class, test$Target)

## [1] 0.5955882

```

We can see accuracy increased by using best CP value

## SVM Model

First model with our selected values of cost and gamma

```

library(e1071)

svmModel<- svm(Target ~ . , data=train, type="C-classification", cost=100)
summary(svmModel)

##
## Call:
## svm(formula = Target ~ ., data = train, type = "C-classification",
##      cost = 100)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##      cost:  100
##
## Number of Support Vectors:  239
##
## ( 112 127 )
##
##
## Number of Classes:  2
##
## Levels:
##  0 1

preds <- predict(svmModel, newdata=test)
preds

##   4   8  10  13  16  28  33  35  39  42  52  55  60  68  95 102 112 113 12
## 6 128
##   1   1   1   1   0   1   1   1   1   1   1   0   1   1   1   0   0   1
## 1   0
## 129 130 138 141 157 164 167 168 180 192 195 196 204 213 214 219 240 258 26

```

```

2 274
## 0 1 1 0 0 0 1 1 1 1 0 0 0 0 1 1 0 1
0 0
## 285 287 300 310 313 339 345 347 356 363 368 389 392 397 399 426 439 458 46
2 463
## 1 0 1 0 1 1 1 0 0 1 0 1 0 1 0 0 0 1
0 0
## 467 471 480 501 503 513 516 517 519 530 538 539 540 544 557 563 566 567 57
5 596
## 0 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0 0
1 0
## 597 598 601 605 606 611 612 628 635 640 645 647 651 656 660 665 667 691 70
4 710
## 0 1 1 1 0 1 0 1 0 0 1 0 1 1 0 1 1 1
1 1
## 721 728 731 738 748 753 757 763 767 775 780 789 791 794 812 814 816 824 83
6 851
## 1 0 1 0 0 0 1 1 0 1 1 0 1 0 1 0 0 0
0 0
## 875 886 891 897 924 926 927 930 951 952 953 958 976 980 986 999
## 1 1 1 0 1 0 1 0 1 1 0 1 1 1 1 0
## Levels: 0 1

table(test$Target, preds)

##      preds
##      0  1
## 0 28 21
## 1 30 57

accuracy(test$Target, preds)

## [1] 0.625

```

type = "C-classification" - binary classification

cost - cost of misclassification:

- if high-> not many misclassified points, margin can be small
- if low -> make more mistakes, margin is larger

Gamma defines how far the influence of a single training example reach.(influence of points either near or far away from the hyperplane.)

- Higher value of gamma → every point has close reach data → chance of overfitting, decision boundary looks wiggly.
- Low value of gamma → every point has far reach data → decision boundary looks smoother

to find best cost and gamma values, use cross validation



This function uses cv to find best value of gamma

Gamma in SVM is usually a value between 0 and 1

```
tunesvm <- tune(svm, Target ~. , data = train, kernel="radial", ranges = list
(gamma = seq(.01, 0.1, by = .01), cost = 100, tunecontrol = tune.control(nrep
eat = 10, sampling = "cross", cross = 10)))
tunesvm

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   gamma cost tunecontrol
##   0.01  100          FALSE
##
## - best performance: 0.3911765
```

*SVM Model using best gamma value obtained by doing CV*

```
SVMmodel_tuned<- svm(Target~., data = train, type="C-classification" , cost =
100, gamma=0.01)
predicted_svm <- predict(SVMmodel_tuned, newdata = test)

accuracy(test$Target, predicted_svm)

## [1] 0.6691176

recall <- function(actual,predictions)
{
  y <- as.vector(table(predictions,actual))
  names(y) <- c("TN","FP","FN","TP")
  recall <- (y["TP"] / (y["TP"]+ y["FN"]))
  return(as.numeric(recall))
}

recall(test$Target, predicted_svm)

## [1] 0.7701149

precision <- function(actual,predictions)
{
  y <- as.vector(table(predictions,actual))
  names(y) <- c("TN","FP","FN","TP")
  precision <- (y["TP"] / (y["TP"]+ y["FP"]))
  return(as.numeric(precision))
}

precision(test$Target, predicted_svm)
```

```
## [1] 0.7282609

accuracy<- function(actual,predictions)
{
  y <- as.vector(table(predictions,actual))
  names(y) <- c("TN","FP","FN","TP")
  accuracy <- (y["TN"] + y["TP"])/ sum(y)
  return(as.numeric(accuracy))
}
accuracy(test$Target, predicted_svm)

## [1] 0.6691176
```

## NAIVE BAYES MODEL

```
naive_model <- naiveBayes(Target ~ ., data= train)
naive_pred_class <- predict(naive_model, test, type="class", laplace =1)
naive_pred_class

## [1] 0 0 1 1 1 1 1 1 1 1 1 0 0 1 1 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0 1 0 1
1 0 1
## [38] 1 1 0 1 0 1 1 1 0 1 1 1 1 0 1 1 1 1 1 0 1 1 1 0 1 1 1 0 1 1 1 1 1 1 1
0 1 1
## [75] 1 1 0 1 1 1 0 1 1 1 1 1 1 1 1 0 1 0 0 0 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1
0 1 1
## [112] 0 1 1 1 0 0 1 1 0 1 0 1 0 1 1 0 1 1 1 0 1 1 1 1 0 0
## Levels: 0 1

naive_pred_prob <- predict(naive_model, test, type="raw")

#confusion matrix
table(naive_pred_class, test$Target, dnn= c("Prediction", "Actual"))

##           Actual
## Prediction 0  1
##           0 19 17
##           1 30 70

accuracy(naive_pred_class, test$Target)

## [1] 0.6544118
```


## Evaluation of Models using ROC Curves

```
library(ROCR)

# 2 main functions: prediction & performance
# prediction(True Labels, Predicted Probabilities for positive class)
```

```
# ROC Curve for SVM model
```

```
svmModel1<- svm(Target ~ . , data=train, probability = TRUE, type="C-classifi  
cation", gamma=0.01, cost=100, decision.values = TRUE)  
pred_svm <- predict(svmModel1, newdata=test, probability = TRUE, decision.val  
ues = TRUE)  
# returns predicted class, and probabilities of belonging on each class  
pred_prob_svm <- attr(pred_svm, "probabilities")  
# store results of the probabilities of being in each class  
pred_prob_svm_good <- pred_prob_svm[,2]  
# store prob of ONLY being in class 1 - Good  
pred <- prediction(pred_prob_svm_good, test$Target)  
  
#ROC CURVE CHART  
perf_roc_svm <- performance(pred, "tpr", "fpr")  
plot(perf_roc_svm)
```




```
#AUC
```

```
auc_svm <- performance(pred, "auc")  
auc_svm <- unlist(slot(auc_svm, "y.values"))  
auc_svm  
  
## [1] 0.7037298
```

```
# ROC Curve for Logistic Regression Model
```

```
pred_lr <- prediction(predictions, test$Target)  
perf_roc_lr <- performance(pred_lr, "tpr", "fpr")  
plot(perf_roc_lr)
```




```
#AUC
auc_LR <- performance(pred_lr, "auc")
auc_LR <- unlist(slot(auc_LR, "y.values"))
auc_LR
```

```
## [1] 0.7447807
```

```
# ROC Curve for Decision Tree Model
```

```
tree_pruned_pred_probs <- predict(tree_pruned, test)
#probability of being in each class
tree_pruned_pred_probs_positive <- tree_pruned_pred_probs[,2]
#probability of being in class 1-Good

pred_decision_tree <- prediction(tree_pruned_pred_probs_positive, test$Target)
perf_roc_decision_tree <- performance(pred_decision_tree, "tpr", "fpr")
plot(perf_roc_decision_tree)
```




```
#AUC
auc_DT <- performance(pred_decision_tree, "auc" )
auc_DT <- unlist(slot(auc_DT, "y.values"))
auc_DT

## [1] 0.575651
```

#### *# ROC Curve for Naive Bayes Model*

```
naive_pred_prob_yes <- naive_pred_prob[,2]
pred_nb <- prediction(naive_pred_prob_yes, test$Target)
perf_roc_nb <- performance(pred_nb, "tpr", "fpr")
plot(perf_roc_nb)
```



```
#AUC
auc_NB <- performance(pred_nb, "auc")
auc_NB <- unlist(slot(auc_NB, "y.values"))
auc_NB

## [1] 0.6572836
```

Overall, our Logistic Regression Model is performing better than SVM, DT, and Naive Bayes.

The accuracy of the Logistic Regression Model: 0.6838235, AUC: 0.7447807

The accuracy of the SVM: 0.6691176 , AUC: 0.7037298

The accuracy of the DT: 0.5955882 , AUC:0.575651

The accuracy of the NB: 0.6544118 , AUC:0.6572836